

MapReduce

Assignment Project Exam Help
<https://tutorcs.com>
WeChat: cstutorcs
CMPSC/DS 410

Reading

- Data Intensive Text Processing with MapReduce, Ch 2.
- Jeffrey Dean and Sanjay Ghemawat. **MapReduce: Simplified Data Processing on Large Clusters**. *Communications of the ACM* January 2008, Vol 52. No.1.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Brief History

- Invented at Google to process large data.
 - Open source implementation developed with help from Yahoo!
 - Named Hadoop
 - Not an acronym for anything
 - Named after toy elephant
- Assignment Project Exam Help
<https://tutorcs.com>



WeChat: cstutorcs

- Because why name it after something descriptive?
- Now an Apache project

Motivation

The Problem

- Large textfile
- Need to write wordcount
 - For each word, count the number of times it occurs
 - File might be too big to fit in memory
 - Number of words might be too big to fit in memory
- Applications:
 - Find word frequencies of English words in news stories
 - Find word frequencies of in Tweets

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Small data solution

```
def wordcount(filename):  
    counter = {}  
    with open(filename) as infile:  
        for line in infile:  
            words = line.split()  
            for w in words:  
                counter[w] = counter.get(w, 0) + 1  
    return counter
```

- Problem 1: file might not fit on your machine
- Problem 2: counter might not fit
- Counter is an associative array.
 - key, value pairs
- How would it be split across multiple machines?
 - Each machine is responsible for a set of keys.

Motivation

Parallel Reads

```
Node j (mapper):  
    Read file shards on Node j.  
    For each line encountered:  
        words = line.split()  
        for w in words:  
            send message (w,1) to reducer  
Reducer Node (does the counting/aggregation):  
    Receive messages, increment word counts
```

- Problem 1: Data not all on one machine.
 - Each shard is 64 MB worth of text lines.
 - Multiple nodes read separate file shards in parallel.
- One node (reducer) stores the dictionary.
- Each mapper sends (many) messages to reducer node.
 - Message key: "word", message value: "1"
 - Reducer interprets message to mean "increment count of word by 1"

Motivation

Parallel Reads

Parallel Aggregation

```
Node j (mapper):
  Read file shards on Node j.
  For each line encountered:
    words = line.split()
    for w in words:
      determine reducer node w
      send message (w,1) to appropriate reducer node
Reducer Node i:
  Receive messages, increment word counts
  At then end, save output to HDFS (as a file shard)
```

- Problem 2: Words cannot fit on one machine.
- Multiple Reducers
 - Each reducer is associated with set of words
 - E.g., Reducer 0 gets words where `python hash(words) % 6 == 0`
 - Reducer 1 gets words where `python hash(words) % 6 == 1`
 - In this specific example there are only 6 reducers (why?)
 - Mapper must send (w, 1) message to appropriate reducer
- Output of reducers cannot fit on one machine.
 - Output is sharded file, shard i produced by reducer i

Motivation

Parallel Reads

Parallel Aggregation

Map Reduce ideas

MapReduce and Hadoop

- You define a:
 - mapper: what key/value messages to create from each line
 - partitioner: decide who is responsible for receiving each key
 - reducer: what to do with the messages received at a node
 - sorter: what order should the messages be received in
 - combiner: an optimization suggestion
- Hadoop:
 - Written in Java
 - MRJob: Convenient Python interface
 - Penalty for convenience:
 - Slower: uses hadoop streaming interface (communicates with java code through stdin and stdout)
 - Less control of useful components (e.g., partitioner)
 - Benefits: shorter code

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Mappers

- **mapper** function
 - inputs a key,value pair.
 - E.g., key = line number, value = line
 - E.g., key = null, value = line
 - outputs 0, 1, or more key value pairs.
- What Hadoop does automatically:
 - Finds Mapper Nodes that have file shards or can get them quickly.
 - Each Mapper Node reads its file shards
 - For every line, it calls the mapper function on it
 - Then it saves the results to local disk (not HDFS)
 - Messages for reducer 1 are saved in 1 file.
 - Sorted by keys assigned to that reducer
 - Messages for reducer 2 are saved in another.
 - Sorted by keys assigned to that reducer
 - All you need to do is write map function

```
from mrjob.job import MRJob

class WordCount(MRJob):
    def mapper(self, key, line):
        words = line.split()
        for w in words:
            yield (w, 1)

    def reducer(self, key, values):
        #TODO
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Partitioners

- Mapper creates key,value pairs
 - e.g., ("it", 1), ("was", 1), ("dark", 1), ("rainy", 1), ("night", 1)
 - keys are the first part of tuple (words)
 - values are the second part (the numbers)
- (key, value) pairs are first saved to file then sent to **reducers** nodes
 - you specify how many reducers
- Partitioner:
 - function that takes a key and outputs a number
 - e.g., partition("was") = 2
 - so ("was", 1) is sent to reducer 2
 - if not specified, Hadoop will use a default
 - custom partitioner poorly supported by MRJob

```
from mrjob.job import MRJob

class WordCount(MRJob):
    def mapper(self, key, line):
        words = line.split()
        for w in words:
            yield w, 1

    def reducer(self, key, values):
```

- What Hadoop does automatically:
 - Reads output from each mapper
 - Collects key, value pairs for each reducer
 - Separate file for each reducer
 - Each file sorted by keys for that reducer
 - Sends (key, value) to appropriate reducer
 - Ensures reducers receive their keys in sorted order
 - "Shuffle and sort" phase
 - you can also specify sort order

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Reducers

- Each reducer is responsible for set of keys
- **reduce** function
 - inputs a key, and list of values
 - values from all messages with that key
 - outputs a key/value pair
- What Hadoop does automatically:
 - Collects incoming key/value pairs
 - Gets these files from multiple mappers
 - Merges them to maintain sorted order
 - Groups by key to create key, value_list
 - Calls reducer(key, value_list) on each one
 - Saves result to HDFS

```
from mrjob.job import MRJob

class WordCount(MRJob):
    def mapper(self, key, line):
        words = line.split()
        for w in words:
            yield (w, 1)

    def reducer(self, key, values):
        yield (key, sum(values))

if __name__ == '__main__':
    WordCount.run()
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Mapper Node 1 input shards:

■ That is that

■ to be or
not to be

Mapper Node 2 input shards:

■ That is big

Mapper 1's function outputs:

■ (That, 1)
(is, 1)
(that, 1)
(to, 1)
(be, 1)
(or, 1)
(not, 1)
(to, 1)
(be, 1)

Mapper 2's function outputs:

■ (That, 1)
(is, 1)
(big, 1)

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Mapper Node 1 outputs:

Sorted file for Reducer 1:

```
(be, 1)
(be, 1)
(not, 1)
(that, 1)
(That, 1)
```

Sorted file for Reducer 2

```
(is, 1)
(or, 1)
(to, 1)
(to, 1)
```

Mapper Node 2
outputs:

Sorted file for
Reducer 1

```
(big, 1)
(That, 1)
```

Sorted file for
Reducer 2:

```
(is, 1)
```

- Note: reducer input is locally sorted, not globally sorted.
- Reducer gets inputs from multiple mappers
 - Merges inputs to maintain sorted order
 - Collects values with same key

Reducer **Node 1** input

```
be, [1, 1]
big, [1]
not, [1]
that, [1]
That, [1, 1]
```

Reducer **Node 2** input

```
is, [1, 1]
or, [1]
to, [1, 1]
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Reducer **Node 1** input

```
be, [1, 1]
big, [1]
not, [1]
that, [1]
That, [1, 1]
```

Reducer Node 1 output (shard 1):

```
be 2
big 1
not 1
that 1
That 2
```

Reducer **Node 2** input

```
is, [1, 1]
or, [1]
to, [1, 1]
```

Reducer Node 2 output (shard 2):

```
is 2
or 1
to 1
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Programming Considerations

- Native language of Hadoop is Java
 - Writing jobs in java provides type safety
 - Writing jobs in java provides more control (e.g., custom partitioner)
- Need to worry about garbage collection.
 - Object churn: occurs when much of time is spent creating/garbage collecting objects.
 - Minimize churn by object re-use.
 - See WordCount2.java <https://tutorcs.com>
 - The Key is created once per mapper (Text word = new Text();)
 - Then reused inside map function (word.set(token))
- Hadoop supports **counters** for maintaining job statistics
 - You control when to increment/decrement counters
 - Counters work in a distributed setting because increment/decrement are associative
 - increment/decrement messages can arrive out of order
 - but order does not matter to final result
 - counter only useful when job has finished

Hadoop Streaming

- Hadoop supports other languages (C, R, Python, etc.)
- Hadoop streaming
 - mapper and reducer written in language of your choice
 - key/value inputs sent to STDIN
 - key/value outputs read in STDOUT
 - tabs, spaces, non-ascii characters and language encodings are important (can cause crashes).
- Slower since requires communicating with/running code outside JVM
- mrjob provides an python interface to jobs using hadoop streaming.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Gotchas


- Multiple mappers run on different nodes.
 - Changing a class variable in one mapper does not affect other nodes!
 - Mappers run independently from each other
 - Why is this good for performance?
 - Why is this good for correctness and reduces potential for bugs?
- Reducers run on different nodes.
- Reducers and mappers only communicate in one way:
 - key,value pairs emitted by mapper are sent to some reducer
 - changing/setting a value in a mapper does not affect reducer

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

YARN Architecture



- Yet Another Resource Negotiator
- Comes with Hadoop 2.0
- Allows Hadoop, Spark, etc. to coexist on cluster.

- Components:
 - Container
 - Node Manager
 - Resource manager
 - App Master


source:

<http://hadoop.apache.org/docs/current/hadoop-yarn/hadoop-yarn-site/YARN.html>

YARN Architecture

Container

- A share of cpu, memory, etc.
- Analogous to virtual machine




<https://tutorcs.com>

WeChat: cstutorcs

YARN Architecture

Node Manager




- Monitors containers
- Reports usage to Resource Manager

<https://tutorcs.com>

WeChat: cstutorcs

YARN Architecture

Resource Manager




- Scheduler and Applications Manager
- Scheduler maintains queue of waiting jobs
- Applications manager
 - accepts new jobs
 - decides first container for a job

<https://tutorcs.com>

WeChat: cstutorcs

YARN Architecture

App Master



- Starts in first container of your job
- Negotiates additional containers from Resource Manager
 - e.g., I need 10 mappers
 - e.g., I need 12 reducers
- Tracks progress of your job
 - e.g., Mapper 1 failed
 - Find new mappers with replicas of data shards Mapper 1 was using
 - Use those mappers to redo that part of computation
 - e.g., Mapper 2 is too slow, restart Mapper 2 in new container