

Advanced MapReduce

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

CMPSC/DS 410

MapReduce Components

- Mapper
 - Can be written in Java
 - Can be written in language of your choosing with hadoop-streaming
- Reducer
 - Can be written in Java
 - Can be written in language of your choosing with hadoop-streaming
- Combiner (optional)
- Partitioner
 - Write it in java
 - Can tell hadoop-streaming (or mrjob) to use it
- Sorter
 - Write it in java

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Remember the WordCount!

```
from mrjob.job import MRJob
```

```
class WordCount(MRJob):
```

```
    def mapper(self, key, line):
```

```
        words = line.split()
```

```
        for w in words:
```

```
            yield (w, 1)
```

```
    def reducer(self, key, values):
```

```
        yield (key, sum(values))
```

```
if __name__ == '__main__':
```

```
    WordCount.run()
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

But we want more

- Suppose we want to know some Statistics:

- total number of words
- number of words starting with A
- number of words starting with B
- etc.

- How would we do it?

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Option 1

- Run a mapreduce job on output of WordCount (using Steps api of MRJOB)

```
def mapper(self, key, value):  
    yield "Total", value  
    first_letter = key[0].upper()  
    yield first_letter, value  
  
def reducer(self, key, value_list):  
    yield key, sum(value_list)
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Option 1

- Run a mapreduce job on output of WordCount (using Steps api of MRJOB)

```
def mapper(self, key, value):  
    yield "Total", value  
    first_letter = key[0].upper()  
    yield first_letter, value  
  
def reducer(self, key, value_list):  
    yield key, sum(value_list)
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

- "With big data comes big responsibility" -- Ben Parker (Spiderman's Uncle)
- This code is wrong! (But why? It gives me the right answer).
 - Inefficient
 - Getting the right answer **slowly**
 - Second mapreduce job is an extra transfer of data
 - Can we get wordcount and statistics in 1 mapreduce job?

Option 2

```
class WordCountStats(MRJob):
    def mapper(self, key, line):
        words = line.split()
        for w in words:
            yield ("_"+w, 1)
        yield "Total_", 1
        first_letter = key[0].upper()
        yield first_letter + "_", 1

    def reducer(self, key, values):
        yield (key, sum(values))
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

- Why the underscore?
- Yucky output. Why?

Option 2

```
class WordCountStats(MRJob):  
    def mapper(self, key, line):  
        words = line.split()  
        for w in words:  
            yield ("_" + w, 1)  
        yield "Total_", 1  
        first_letter = key[0].upper()  
        yield first_letter + "_", 1  
  
    def reducer(self, key, values):  
        yield (key, sum(values))
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

- Where is the total? Where are the A words?
 - Scattered on different part files (shards)
 - Within a part file, mixed in with the normal word count information
 - How do we fix this?

Wait until someone says "Partitioner"

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Option 3

```
from mrjob.job import MRJob

class WordCountStats(MRJob):
    def mapper(self, key, line):
        words = line.split()
        for w in words:
            yield "_" + w, 1
        yield "Total_", 1
        first_letter = key[0].upper()
        yield first_letter + "_", 1

    def reducer(self, key, values):
        yield (key, sum(values))

if __name__ == '__main__':
    WordCountStats.run()
```

```
int getPartition(Key key, Value value, int numPart) {
    if key.startsWith("_") { //normal word
        partition = 1 + (key.hash() % (numPart - 1))
    }
    else { //Summary
        partition = 0
    }
    return partition
}
```

Assignment Project Exam Help

<https://tutores.com>

WeChat: cstutores

- Statistics go to first partition
- Wordcounts distributed across the rest of the partitions
- Good?

Big Data

- Big data can easily overwhelm big resources
- Have to use resources efficiently
- In this sense the WordCountStats program is still wrong.
- To see why, let us look at LetterCount

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Counting letters

```
class LetterCount(MRJob):  
    def mapper(self, key, line):  
        for symbol in line:  
            if isletter(symbol):  
                yield (symbol, 1)  
  
    def reducer(self, key, values):  
        yield (key, sum(values))
```

Assignment Project Exam Help

<https://tutorcs.com>

- What is the size of the data being shipped from mappers to reducers?

WeChat: cstutorcs

Counting letters

```
class LetterCount(MRJob):  
    def mapper(self, key, line):  
        for symbol in line:  
            if isletter(symbol):  
                yield (symbol, 1)  
  
    def reducer(self, key, values):  
        yield (key, sum(values))
```

Assignment Project Exam Help

<https://tutorcs.com>

- What is the size of the data being shipped from mappers to reducers?
 - L number of characters in input
 - B is size of a character
 - N is size of an integer
 - Roughly $L(B+N)$ being shipped across network
 - Could be larger than original input file
 - What to do?

WeChat: cstutorcs

Combining

```
class LetterCount(MRJob):  
    def mapper(self, key, line):  
        for symbol in line:  
            if isletter(symbol):  
                yield (symbol, 1)  
  
    def reducer(self, key, values):  
        yield (key, sum(values))
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

- Why wait until reducer to start adding up values?
- We can **combine** messages with the same key by adding up their values.

Setup and Teardown functions

```
def mapper_init(self):  
    # initialization code  
  
def mapper_final(self):  
    # cleanup code  
  
def reducer_init(self):  
    # initialization code  
  
def reducer_final(self):  
    # cleanup code
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

- You can override them
- They are called automatically by mrjob
- Call sequence for each mapper:
 - first `mapper_init()` is called
 - then for every line, `mapper(key, value)` is called
 - then `mapper_final()` is called
 - similar sequence for reducers

Setup and Teardown

- `mapper_init()`
 - can initialize class variables (`self.cache={}`)
 - these variables are available when the **same** mapper processes subsequent lines
 - i.e. a mapper has its own state
 - it can modify its own state
 - it cannot modify the state of other mappers (they are on different machines and may run at different times)
 - can output key,value pairs (using **yield**)
- `mapper_final()`
 - cleans up at the end
 - can output key, value pairs
 - can close any files that were opened

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

How to use for letter count

```
def mapper(self, key, line):  
    for symbol in line:  
        if isletter(symbol):  
            yield (symbol, 1)  
  
def reducer(self, key, values):  
    yield (key, sum(values))
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Hint:

```
def mapper(self, key, line):
    for symbol in line:
        if isletter(symbol):
            yield (symbol, 1)

def reducer(self, key, values):
    yield (key, sum(values))
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

- mapper_init() initializes a dictionary (key=letter, value=how many times seen so far)
- mapper() updates the counts, yields nothing but sends heartbeat messages periodically
 - call **set_status(message)** to send heartbeat message
 - why is it needed? what happens if heartbeat message is not sent?
- mapper_final() outputs key value pairs from the dictionary
- Combining messages inside mapper is called **in memory combining**
- Combiners in hadoop also specify how to combine messages
 - But there are no guarantees on when/if/how-many-times a combiner will run
 - Hence in memory combining is preferred
- Why should we use only 1 reducer for this specific problem?

So why is this wrong?

```
from mrjob.job import MRJob

class WordCountStats(MRJob):
    def mapper(self, key, line):
        words = line.split()
        for w in words:
            yield ("_" + w, 1)
        yield "Total_", 1
        first_letter = key[0].upper()
        yield first_letter + "_", 1

    def reducer(self, key, values):
        yield (key, sum(values))

if __name__ == '__main__':
    WordCountStats.run()
```

```
int getPartition(Key key, Value value, int numPart) {
    if key.startsWith("_") { //normal word
        partition = 1 + (key.hash() % (numPart - 1))
    }
    else { //Summary
        partition = 0
    }
    return partition
}
```

Assignment Project Exam Help

<https://tutores.com>

WeChat: cstutorcs

What was wrong

- Reducer 0 gets message "Total_", 1 for **every** word in input
 - Assumes 1 reducer can store something as big as the entire input
 - That is small data programming
 - For big files your job will crash
- No in memory combining

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

How to add in memory combining?

```
from mrjob.job import MRJob

class WordCountStats(MRJob):
    def mapper(self, key, line):
        words = line.split()
        for w in words:
            yield "_" + w, 1
        yield "Total_", 1
        first_letter = key[0].upper()
        yield first_letter + "_", 1

    def reducer(self, key, values):
        yield (key, sum(values))

if __name__ == '__main__':
    WordCountStats.run()
```

```
int getPartition(Key key, Value value, int numPart) {
    if key.startsWith("_") { //normal word
        partition = 1 + (key.hash() % (numPart - 1))
    } else { //Summary
        partition = 0
    }
    return partition
}
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Note:

- Don't let the dictionary (or other data) stored by mapper get too big
 - Otherwise it might not fit in memory
 - Avoid small data solutions!
 - Flush the cache
 - output key value pairs when the cache gets too big
 - then clear the cache

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

What we now know:

- Using partitioners to send messages where we want
- Using in memory combining to reduce the number of messages sent
 - Reduces network traffic
 - Makes the Shuffle and Sort phase faster (less data to shuffle and sort)
 - Helps avoid overloading a reduce with too many messages

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Next problem: Word rates

- WordCount output is **word, count**
- WordRates output is **word, rate**
 - $\text{rate} = (\text{number of times word appears}) / (\text{total number of words in input})$
- How to do it?

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

General Strategy

- Mapper emits two kinds of messages:
 - word, count
 - Total_0, count
 - emitted by mapper_final()
 - represents total number of words seen by this mapper
 - this message is sent to Reducer 0 (via partitioner)
 - Total_1, count
 - emitted by mapper_final()
 - represents total number of words seen by this mapper
 - this message is sent to Reducer 1 (via partitioner)
 - etc.
- Sorter
 - makes sure Total_0 appears before all normal words
 - So Reducer 0 will receive key=Total\0 and value_list as very first key, value_list pair
 - Reducer sums them up, now knows total and can generate word rate for every word it sees
 - make sure Total_1 appears before all normal words
 - So Reducer 1 can do its job.
 - etc.
- Sorter + Partitioner in java can control delivery destination and order

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Sorter

- Write a sorter for word rate

```
int compare(WritableComparable w1, WritableComparable w2) {  
    //return 1 if w1 is "greater" than w2  
    //return 0 if they are equal  
    //return -1 if w1 is "less" than w2  
}
```

Assignment Project Exam Help

<https://tutorcs.com>

- Sorter requirements:
 - sorter works on keys
 - only return 0 if the keys are the same!
 - antisymmetric: $\text{compare}(w1, w2) == -\text{compare}(w2, w1)$
 - transitive property:
 - if $\text{compare}(w1, w2) == 1$
 - and $\text{compare}(w2, w3) == 1$
 - then $\text{compare}(w1, w3) == 1$

WeChat: cstutorcs

Exercise

- Use mapreduce to output word,count
- but only for words that appear more frequently than "like"

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Joins

- File 1 (Customers):
 - key = None
 - value = customer id, city
- File 2 (Orders):
 - key = None
 - value = customer id, item, amount
- Need to do a join
 - Find the total amount of purchases for each city
 - `SELECT city, SUM(amount) FROM Customers C, Orders O WHERE C.id=O.id`
 - Can specify multiple input files for mrjob (just list them)

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Joins

```
def mapper(self, key, value):  
    parts = value.split(",")  
    if len(parts) == 2: # Customers table  
        id = parts[0]  
        city = parts[1]  
        yield id, (city, "Customers")  
    else: # Orders table  
        id = parts[0]  
        item = parts[1]  
        amount = parts[2]  
        yield id, (amount, "Orders")
```

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

- Joining Customers and Orders on id
- Mapper ensures co-location of file
 - for both files, id is the key
 - partitioner decides which reducers get assigned which id
 - records from both files with same id go to same reducer (co-location)
 - records tagged with info about file they belong to

Joins

```
def reducer(self, key, value_list):  
    amount = 0  
    for (a,b) in value_list:  
        if b == "Customers":  
            city = a  
        else:  
            amount = amount + int(a)  
    yield city, amount
```

Assignment Project Exam Help

<https://tutorcs.com>

- Reducer finishes the join
 - Here we assume a customer does not appear multiple times in Customers table
 - In practice we will need to check this.
 - How do we check? What to do if check fails?
 - What if customer has no orders?
 - What if order has invalid customer id (no customer exists?)
- Are we done?
- How do we make it faster?

WeChat: cstutorcs

Joins

```
def reducer(self, key, value_list):  
    amount = 0  
    for (a,b) in value_list:  
        if b == "Customers":  
            city = a  
        else:  
            amount = amount + int(a)  
    yield city, amount
```

Assignment Project Exam Help

<https://tutorcs.com>

- Are we done?
 - No
 - Second mapreduce step must add amounts for same city
- How do we make it faster?
 - in-memory combining inside reducer
- Main point: mapper co-locates records from two files by using join value as the key

WeChat: cstutorcs

PageRank

- One of the initial motivations for mapreduce.
- Every web page (or site) has a reputation score.
- Every web page has outgoing links
- Initially every page has the same score (1).
- Input file:
 - key = website
 - value = (score, list of websites)
- Algorithm operates in rounds. In each round
 - A site takes its score,
 - divides it evenly among outgoing links
 - the new score of a page is sum of scores it receives
- Original algorithm has random restarts which we ignore

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

PageRank

```
def mapper(key, value):  
    score = value[0]  
    neighbors = value[1]  
    yield key, ("neighbors", neighbors) #why?  
    for site in neighbors:  
        yield site, ("score", score/len(neighbors))
```

- mapper sends the neighborlist
- mapper sends a share of the score to each neighbor

<https://tutorcs.com>

WeChat: cstutorcs

PageRank

```
def mapper(key, value):  
    score = value[0]  
    neighbors = value[1]  
    yield key, ("neighbors", neighbors) #why?  
    for site in neighbors:  
        yield site, ("score", score/len(neighbors))
```

```
def reducer(key, value_list):  
    total = 0  
    for v in value_list:  
        if v[0] == "neighbors":  
            neighbor_list = v[0]  
        else:  
            total = total + int(v[1])  
    yield key, (total, neighbor_list)
```

<https://tutorcs.com>

WeChat: cstutorcs

- In reducer: one of the values will be neighbor list
- The other values will be scores received
- In actual code:
 - in mrjob, each value is a string
 - in java it is more flexible